

```
FIRST_NAME = "Mel"
```

```
LAST_NAME = "Gerst"
```

```
STUDENT_ID = "800995291"
```

```
Final Class Project - ITCS 5154
```

```
Dog Breed Classifier - Student Mel Gerst
```

```
Duplicating project originally by TechVidvan
```

Resources ×



You are not subscribed. [Learn more](#)

Available: 50.66 compute units

Usage rate: approximately 10.59 per hour

You have 1 active session.

[Manage sessions](#)

Want more memory and disk space? ×

[Upgrade to Colab Pro](#)

Python 3 Google Compute Engine backend
(GPU)

Showing resources from 4:40 PM to 6:15 PM


System RAM
17.6 / 83.5 GB



GPU RAM
19.7 / 40.0 GB



Disk
35.5 / 112.6 GB




```
# Dog Breed Classifier
# ITCS 5154 – Student Mel Gerst
# Duplicating project by TechVidvan
# Import necessary packages for dog breed
import cv2
import numpy as np
import pandas as pd
import tensorflow
import pathlib
import os

from tensorflow.keras.preprocessing.image
from sklearn.model_selection import train_
from sklearn.preprocessing import LabelEnc
from tensorflow.keras.models import load_n
from tensorflow.keras.optimizers import RM
from tensorflow.keras.layers import Dense,
from tensorflow.keras.applications.resnet_

print("Imports Complete")
print(pathlib.Path().resolve())

from google.colab import drive
# drive._mount('/content/drive')
drive.mount('/content/drive/')
# My files are mounted in Google drive for
# file_path = '/content/drive/My Drive/dat
```

 Imports Complete
/content
Drive already mounted at /content/driv

```
# Initialize Variables
encoder = LabelEncoder()
image_size = 224
breed_count = 60
batch_size = 64

# Grab input files and data
df_labels = pd.read_csv("/content/drive/My
#store training and testing images folder
training_data = '/content/drive/My Drive/(
testing_data =  '/content/drive/My Drive/(

# Check and print the total number of unique
print("Total number of unique Dog Breeds in d
print(os.listdir(testing_data))
```

➞ Total number of unique Dog Breeds in d
['06b727fc8e24e46fd7ea78b08091cab5.jpg

```
# Drop breeds considered to 60 breeds to s
breed_dict = list(df_labels['breed'].value
new_list = sorted(breed_dict,reverse=True)
# Limit dataset to have only those 60 unic
df_labels = df_labels.query('breed in @nev
# Add new column which will contain image
df_labels['img_file'] = df_labels['id'].ap
print("Total number of unique Dog Breeds u
print("The breeds used for training and te

# Create a numpy array of the shape (numbe
# Input for model
train_x = np.zeros((len(df_labels), image_
```

➞ Total number of unique Dog Breeds used
The breeds used for training and testi
<ipython-input-3-e80864a923a5>:7: Sett
A value is trying to be set on a copy
Try using .loc[row_indexer,col_indexer

See the caveats in the documentation:
df_labels['img_file'] = df_labels['i

```
# #iterate over img_file column of our data
# for i, img_id in enumerate(df_labels['img_file']):
#     # Read the image file and convert into array
#     # Resize all images to one dimension (224,224)
#     # We will get array with the shape of (224,224,3) where 3 is the RGB channels
#     img = cv2.resize(cv2.imread(training_images[i]), (224, 224))
#     # Scale array into the range of -1 to 1
#     # Preprocess the array and expand its dimension
#     img_array = preprocess_input(np.expand_dims(img, axis=-1))
#     # Update the train_x variable with new array
#     train_x[i] = img_array
```

```
# print(train_x.shape)
# np.save('/content/drive/My Drive/ColabNotebooks/dog_breed_classifier/train_x.npy', train_x)
```

```
train_x = np.load('/content/drive/My Drive/ColabNotebooks/dog_breed_classifier/train_x.npy')
print(train_x.shape)
```

 (5175, 224, 224, 3)

```
# This will be target for model.
# Convert breed names into numerical format
train_y = encoder.fit_transform(df_labels['breed'])
```

```
# Split the dataset in the ratio of 80:20.
#80% for training and 20% for testing purpose
x_train, x_test, y_train, y_test = train_test_split(train_x, train_y, test_size=0.2, random_state=42)
```

```
#Image augmentation using ImageDataGenerator
train_datagen = ImageDataGenerator(rotation_range=30,
                                   width_shift_range=0.1,
                                   height_shift_range=0.1,
                                   shear_range=0.1,
                                   zoom_range=0.1,
                                   horizontal_flip=True,
                                   fill_mode='nearest')

# Generate images for training sets
train_generator = train_datagen.flow(x_train, y_train,
                                     batch_size=batch_size)

# Same process for Testing sets also by default
test_datagen = ImageDataGenerator()

test_generator = test_datagen.flow(x_test, y_test,
                                   batch_size=batch_size)
```

```

# Model #1 - Build the model using ResNet50
# Weights for our network will be from of
# We will not include the first Dense layer
resnet = ResNet50V2(input_shape = [image_size, image_size, 3])
# Freeze all trainable layers and train on
for layer in resnet.layers:
    layer.trainable = False

# Add global average pooling layer and Batch Normalization
x = resnet.output
x = BatchNormalization()(x)
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
# Add fully connected layer
x = Dense(1024, activation='relu')(x)
x = Dropout(0.5)(x)

# Add output layer having the shape equal to the number of breeds
predictions = Dense(breed_count, activation='softmax')(x)

# Create model class with inputs and outputs
model = Model(inputs=resnet.input, outputs=predictions)

```

```

# model.summary()

# Set the num_epochs for model training and learning rate
num_epochs = 20
learning_rate = 1e-3

# Using RMSprop optimizer compile or build the model
optimizer = RMSprop(learning_rate=learning_rate)
model.compile(optimizer=optimizer,
              loss='sparse_categorical_crossentropy',
              metrics=["accuracy"])

# Fit the training generator data and train the model
model.fit(train_generator,
          steps_per_epoch= x_train.steps_per_epoch,
          epochs= num_epochs,
          validation_data= test_generator,
          validation_steps= x_test.steps_per_epoch)

test_loss, test_accuracy = model.evaluate(test_data_loader)

```

```
print(f'Resnet Test accuracy: {test_accuara

# Save the model for prediction
model.save("model.keras")
```

```

Epoch 1/20
/usr/local/lib/python3.10/dist-package
self._warn_if_super_not_called()
64/64 ██████████ 69s 789ms/s
Epoch 2/20
1/64 ██████████ 1s 27ms/st
self.gen.throw(typ, value, traceback
64/64 ██████████ 5s 83ms/st
Epoch 3/20
64/64 ██████████ 42s 583ms/s
Epoch 4/20
64/64 ██████████ 0s 405us/st
Epoch 5/20
64/64 ██████████ 42s 585ms/s
Epoch 6/20
64/64 ██████████ 0s 408us/st
Epoch 7/20
64/64 ██████████ 42s 583ms/s
Epoch 8/20
64/64 ██████████ 0s 394us/st
Epoch 9/20
64/64 ██████████ 42s 583ms/s
Epoch 10/20
64/64 ██████████ 0s 400us/st
Epoch 11/20
64/64 ██████████ 42s 582ms/s
Epoch 12/20
64/64 ██████████ 0s 417us/st
Epoch 13/20
64/64 ██████████ 42s 580ms/s
Epoch 14/20
64/64 ██████████ 0s 400us/st
Epoch 15/20
64/64 ██████████ 42s 581ms/s
Epoch 16/20
64/64 ██████████ 0s 404us/st
Epoch 17/20
64/64 ██████████ 42s 581ms/s
Epoch 18/20
64/64 ██████████ 0s 400us/st
Epoch 19/20
64/64 ██████████ 42s 581ms/s
Epoch 20/20
64/64 ██████████ 0s 397us/st
33/33 ██████████ 7s 15ms/st
Resnet Test accuracy: 0.78

```



```
# Load the model
model = load_model("model.keras")

# Get the image of the dog #1 for prediction
pred_img_path = '/content/drive/My Drive/Colab Notebooks/DogBreedClassifier/1001ThysfOEmHutnf-6g_kia8--_kdP-n?authuser=1#scrollTo=avIJe5JpzocF'
# Read the image file and convert into numpy array
# Resize all images to one dimension i.e. 224x224
pred_img_array = cv2.resize(cv2.imread(pred_img_path), (224, 224))
# Scale array into the range of -1 to 1.
# Expand the dimension on the axis 0 and now it's 1x224x224x3
pred_img_array = preprocess_input(np.expand_dims(pred_img_array, axis=0))

# Feed the model with the image array for prediction
pred_val = model.predict(np.array(pred_img_array))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path), (224, 224)))
# Display the predicted breed of dog
predicted_breed = sorted(new_list)[np.argmax(pred_val)]
print("Predicted Breed for this Dog is :", predicted_breed)

# Get the image of the dog #2 for prediction
pred_img_path2 = '/content/drive/My Drive/Colab Notebooks/DogBreedClassifier/1001ThysfOEmHutnf-6g_kia8--_kdP-n?authuser=1#scrollTo=avIJe5JpzocF'
# Read the image file and convert into numpy array
# Resize all images to one dimension i.e. 224x224
pred_img_array2 = cv2.resize(cv2.imread(pred_img_path2), (224, 224))
# Scale array into the range of -1 to 1.
# Expand the dimension on the axis 0 and now it's 1x224x224x3
pred_img_array2 = preprocess_input(np.expand_dims(pred_img_array2, axis=0))

# Feed the model with the image array for prediction
pred_val2 = model.predict(np.array(pred_img_array2))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path2), (224, 224)))
# Display the predicted breed of dog
predicted_breed2 = sorted(new_list)[np.argmax(pred_val2)]
print("Predicted Breed for this Dog is :", predicted_breed2)

# Get the image of the dog #3 for prediction
pred_img_path3 = '/content/drive/My Drive/Colab Notebooks/DogBreedClassifier/1001ThysfOEmHutnf-6g_kia8--_kdP-n?authuser=1#scrollTo=avIJe5JpzocF'
# Read the image file and convert into numpy array
# Resize all images to one dimension i.e. 224x224
```

```

pred_img_array3 = cv2.resize(cv2.imread(pr
# Scale array into the range of -1 to 1.
# Expand the dimesion on the axis 0 and no
pred_img_array3 = preprocess_input(np.expa

# Feed the model with the image array for
pred_val3 = model.predict(np.array(pred_in
# Display the image of dog
from google.colab.patches import cv2_imshc
cv2_imshow(cv2.resize(cv2.imread(pred_img_
# Display the predicted breed of dog
predicted_breed3 = sorted(new_list)[np.arg
print("Predicted Breed for this Dog is :",

# Get the image of the dog #4 for predicti
pred_img_path4 = '/content/drive/My Drive/
# Read the image file and convert into num
# Resize all images to one dimension i.e.
pred_img_array4 = cv2.resize(cv2.imread(pr
# Scale array into the range of -1 to 1.
# Expand the dimesion on the axis 0 and no
pred_img_array4 = preprocess_input(np.expa

# Feed the model with the image array for
pred_val4 = model.predict(np.array(pred_in
# Display the image of dog
from google.colab.patches import cv2_imshc
cv2_imshow(cv2.resize(cv2.imread(pred_img_
# Display the predicted breed of dog
predicted_breed4 = sorted(new_list)[np.arg
print("Predicted Breed for this Dog is :",

# Get the image of the dog #5 for predicti
pred_img_path5 = '/content/drive/My Drive/
# Read the image file and convert into num
# Resize all images to one dimension i.e.
pred_img_array5 = cv2.resize(cv2.imread(pr
# Scale array into the range of -1 to 1.
# Expand the dimesion on the axis 0 and no
pred_img_array5 = preprocess_input(np.expa

# Feed the model with the image array for
pred_val5 = model.predict(np.array(pred_in
# Display the image of dog
from google.colab.patches import cv2_imshc

```

```
cv2.imshow(cv2.resize(cv2.imread(pred_img_
#Display the predicted breed of dog
predicted_breed5 = sorted(new_list)[np.arg
print("Predicted Breed for this Dog is :",
print("Check image size: ", image_size)
```

1/1 5s 5s/step




Predicted Breed for this Dog is : rott
1/1 0s 24ms/step




Predicted Breed for this Dog is : mini
1/1 0s 25ms/step



Predicted Breed for this Dog is : labr
1/1  0s 24ms/step



Predicted Breed for this Dog is : scot
1/1  0s 24ms/step



Predicted Breed for this Dog is : whip
Check image size: 224

```
#VGG model
import tensorflow as tf
from tensorflow.keras import layers, model
from sklearn.model_selection import train_

# This will be target for model.
# Convert breed names into numerical format
train_y = encoder.fit_transform(df_labels|

# Assuming train_y is one-hot encoded, if
train_y = tf.keras.utils.to_categorical(tr

# Split the data
```

```

x_train, x_test, y_train, y_test = train_t

def create_vgg_model(input_shape, num_classes):
    base_model = tf.keras.applications.VGG16(input_shape, include_top=False)

    # Freeze the base model
    base_model.trainable = False

    model = models.Sequential([
        base_model,
        layers.Flatten(),
        layers.Dense(512, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation='softmax')
    ])

    return model, base_model # Return both

input_shape = (224, 224, 3)
num_classes = 60
model, base_model = create_vgg_model(input_shape, num_classes)

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   epochs=50,
                   batch_size=32,
                   validation_data=(x_test, y_test))

test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_accuracy:.2f}')

base_model.trainable = True # Now base_model is trainable
for layer in base_model.layers[:-4]: # Unfreeze the first 4 layers
    layer.trainable = False

model.compile(optimizer=tf.keras.optimizers.Adam(),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history_finetune = model.fit(x_train, y_train,
                           epochs=10,

```











```
batch_size=32
validation_c
```

```
test_loss, test_accuracy = model.evaluate(
print(f'VGG Test accuracy: {test_accuracy:0.2f}')
```

```

130/130 ————— 3s 23ms/s
Epoch 33/50
130/130 ————— 3s 23ms/s
Epoch 34/50
130/130 ————— 3s 23ms/s
Epoch 35/50
130/130 ————— 3s 23ms/s
Epoch 36/50
130/130 ————— 3s 23ms/s
Epoch 37/50
130/130 ————— 3s 23ms/s
Epoch 38/50
130/130 ————— 3s 23ms/s
Epoch 39/50
130/130 ————— 3s 23ms/s
Epoch 40/50
130/130 ————— 3s 23ms/s
Epoch 41/50
130/130 ————— 3s 23ms/s
Epoch 42/50
130/130 ————— 3s 23ms/s
Epoch 43/50
130/130 ————— 3s 23ms/s
Epoch 44/50
130/130 ————— 3s 23ms/s
Epoch 45/50
130/130 ————— 3s 23ms/s
Epoch 46/50
130/130 ————— 3s 23ms/s
Epoch 47/50
130/130 ————— 3s 23ms/s
Epoch 48/50
130/130 ————— 3s 23ms/s
Epoch 49/50
130/130 ————— 3s 23ms/s
Epoch 50/50
130/130 ————— 3s 23ms/s
33/33 ————— 1s 18ms/ste
Test accuracy: 0.34
Epoch 1/10
130/130 ————— 11s 59ms/
Epoch 2/10
```

```

130/130  3s 26ms/s
Epoch 3/10
130/130  3s 25ms/s
Epoch 4/10
130/130  3s 26ms/s
Epoch 5/10
130/130  3s 26ms/s
Epoch 6/10
130/130  3s 25ms/s
Epoch 7/10
130/130  3s 26ms/s
Epoch 8/10
130/130  3s 26ms/s
Epoch 9/10
130/130  3s 26ms/s
Epoch 10/10
130/130  3s 26ms/s
33/33  1s 18ms/ste

```

```

# Second round of predictions with VGG model
print("Check image size: ", image_size)
# Get the image of the dog #1 for prediction
pred_img_path = '/content/drive/My Drive/Colab Notebooks/DogBreedClassifier/1001ThysfOEmHutnf-6g_kia8--_kdP-n?authuser=1#scrollTo=avIJe5JpzocF'
# Read the image file and convert into numpy array
# Resize all images to one dimension i.e. 224x224
pred_img_array = cv2.resize(cv2.imread(pred_img_path), (image_size, image_size))
# Scale array into the range of -1 to 1.
# Expand the dimension on the axis 0 and not on the axis 1
pred_img_array = preprocess_input(np.expand_dims(pred_img_array, axis=0))

```

```

# Feed the model with the image array for prediction
pred_val = model.predict(np.array(pred_img_array))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path), (image_size, image_size)))
# Display the predicted breed of dog
predicted_breed = sorted(new_list)[np.argmax(pred_val)]
print("Predicted Breed for this Dog is :", predicted_breed)

```

```

# Get the image of the dog #2 for prediction
pred_img_path2 = '/content/drive/My Drive/Colab Notebooks/DogBreedClassifier/1001ThysfOEmHutnf-6g_kia8--_kdP-n?authuser=1#scrollTo=avIJe5JpzocF'
# Read the image file and convert into numpy array
# Resize all images to one dimension i.e. 224x224
pred_img_array2 = cv2.resize(cv2.imread(pred_img_path2), (image_size, image_size))
# Scale array into the range of -1 to 1.
# Expand the dimension on the axis 0 and not on the axis 1
pred_img_array2 = preprocess_input(np.expand_dims(pred_img_array2, axis=0))

```

```

pred_img_array2 = preprocess_input(np.expand_dims(pred_img_array2, axis=0))

# Feed the model with the image array for prediction
pred_val2 = model.predict(np.array(pred_img_array2))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path2), (224, 224)))
# Display the predicted breed of dog
predicted_breed2 = sorted(new_list)[np.argmax(pred_val2)]
print("Predicted Breed for this Dog is :", predicted_breed2)

# Get the image of the dog #3 for prediction
pred_img_path3 = '/content/drive/My Drive/dog_images/dog3.jpg'
# Read the image file and convert into numpy array
img3 = cv2.imread(pred_img_path3)
# Resize all images to one dimension i.e. 224x224
pred_img_array3 = cv2.resize(img3, (224, 224))
# Scale array into the range of -1 to 1.
pred_img_array3 = pred_img_array3 / 255
# Expand the dimension on the axis 0 and convert to float32
pred_img_array3 = preprocess_input(np.expand_dims(pred_img_array3, axis=0))

# Feed the model with the image array for prediction
pred_val3 = model.predict(np.array(pred_img_array3))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path3), (224, 224)))
# Display the predicted breed of dog
predicted_breed3 = sorted(new_list)[np.argmax(pred_val3)]
print("Predicted Breed for this Dog is :", predicted_breed3)

# Get the image of the dog #4 for prediction
pred_img_path4 = '/content/drive/My Drive/dog_images/dog4.jpg'
# Read the image file and convert into numpy array
img4 = cv2.imread(pred_img_path4)
# Resize all images to one dimension i.e. 224x224
pred_img_array4 = cv2.resize(img4, (224, 224))
# Scale array into the range of -1 to 1.
pred_img_array4 = pred_img_array4 / 255
# Expand the dimension on the axis 0 and convert to float32
pred_img_array4 = preprocess_input(np.expand_dims(pred_img_array4, axis=0))

# Feed the model with the image array for prediction
pred_val4 = model.predict(np.array(pred_img_array4))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path4), (224, 224)))
# Display the predicted breed of dog
predicted_breed4 = sorted(new_list)[np.argmax(pred_val4)]
print("Predicted Breed for this Dog is :", predicted_breed4)

```



```

print("Predicted Breed for this Dog is :",

# Get the image of the dog #5 for prediction
pred_img_path5 = '/content/drive/My Drive/
# Read the image file and convert into numpy array
# Resize all images to one dimension i.e. 224x224
pred_img_array5 = cv2.resize(cv2.imread(pred_img_path5), (224, 224))
# Scale array into the range of -1 to 1.
# Expand the dimension on the axis 0 and no. of images
pred_img_array5 = preprocess_input(np.expand_dims(pred_img_array5, axis=0))

# Feed the model with the image array for prediction
pred_val5 = model.predict(np.array(pred_img_array5))
# Display the image of dog
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_path5), (224, 224)))
#Display the predicted breed of dog
predicted_breed5 = sorted(new_list)[np.argmax(pred_val5)]
print("Predicted Breed for this Dog is :",

```




Check image size: 224


1/1  1s 1s/step




Predicted Breed for this Dog is : rott


1/1  0s 20ms/step






Predicted Breed for this Dog is : mini
1/1  0s 21ms/step



Predicted Breed for this Dog is : boxe
1/1  0s 20ms/step



Predicted Breed for this Dog is : iris
1/1  0s 21ms/step



Predicted Breed for this Dog is : grea

```
print(len(np.unique(y_train)))
```

↔ 2

```
# new SimpleCNN model
import torch
import torch.nn as nn
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, Dataset
from sklearn.model_selection import train_test_split
import numpy as np

# Assume train_x and train_y are defined
# train_x: numpy array of shape (5175, 224, 224, 3)
# train_y: numpy array of shape (5175,)

# Split the dataset
# x_train, x_test, y_train, y_test = train_test_split(
#     train_x, train_y, test_size=0.1, random_state=42)

# Set batch size
# batch_size = 64

# Define data augmentation and normalization
train_transforms = transforms.Compose([
    transforms.ToPILImage(), # Convert numpy array to PIL image
    transforms.RandomRotation(45),
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
    transforms.RandomResizedCrop(size=(224, 224)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
])

test_transforms = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
])

# Custom Dataset Class
class CustomDataset(Dataset):
    def __init__(self, images, labels, transform):
        self.images = images
        self.labels = labels
        self.transform = transform
```

```

def __len__(self):
    return len(self.images)

def __getitem__(self, idx):
    image = self.images[idx]
    label = self.labels[idx]

    if self.transform:
        image = self.transform(image)

    return image, label

# Create datasets
y_train_indices = np.argmax(y_train, axis=1)
y_test_indices = np.argmax(y_test, axis=1)
train_dataset = CustomDataset(x_train, y_train_indices)
test_dataset = CustomDataset(x_test, y_test_indices)

# Create DataLoaders
train_loader = DataLoader(train_dataset, batch_size=16)
test_loader = DataLoader(test_dataset, batch_size=16)

# Define the SimpleCNN model
class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)

        # Calculate the input size for the fully connected layers
        self.fc1_input_size = 128 * (224 // 2 // 2) ** 2
        self.fc1 = nn.Linear(self.fc1_input_size, 256)
        self.fc2 = nn.Linear(256, 60) # Output layer

    def forward(self, x):
        x = self.conv1(x)
        x = nn.ReLU()(x)
        x = self.pool(x)

        x = self.conv2(x)
        x = nn.ReLU()(x)
        x = self.pool(x)

```

```

        x = self.conv3(x)
        x = nn.ReLU()(x)
        x = self.pool(x)

        x = torch.flatten(x, 1) # Flatten
        x = self.fc1(x)
        x = nn.ReLU()(x)
        x = self.fc2(x)
        return nn.LogSoftmax(dim=1)(x)

# Initialize model, loss function, and optimizer
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = SimpleCNN().to(device)
criterion = nn.NLLLoss()
optimizer = torch.optim.Adam(model.parameters())

# Training Loop
num_epochs = 200
for epoch in range(num_epochs):
    model.train() # Set model to training
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad() # Clear gradients
        outputs = model(images) # Forward pass
        loss = criterion(outputs, labels)
        loss.backward() # Backward pass
        optimizer.step() # Update weights

    print(f'Epoch [{epoch + 1}/{num_epochs}]')

# Evaluation on the test set
model.eval() # Set model to evaluation mode
test_loss = 0
correct = 0

with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        test_loss += criterion(outputs, labels)
        pred = outputs.argmax(dim=1) # Get predicted class
        correct += (pred == labels).sum().item()

```

```
test_accuracy = correct / len(test_dataset)
print(f'Test Loss: {test_loss/len(test_load
```

```

Epoch [142/200], Loss: 2.4036
Epoch [143/200], Loss: 2.2013
Epoch [144/200], Loss: 2.0762
Epoch [145/200], Loss: 2.0533
Epoch [146/200], Loss: 1.9337
Epoch [147/200], Loss: 2.3631
Epoch [148/200], Loss: 2.1085
Epoch [149/200], Loss: 1.9786
Epoch [150/200], Loss: 2.1217
Epoch [151/200], Loss: 2.1622
Epoch [152/200], Loss: 2.2347
Epoch [153/200], Loss: 2.1163
Epoch [154/200], Loss: 2.4530
Epoch [155/200], Loss: 2.5614
Epoch [156/200], Loss: 2.2602
Epoch [157/200], Loss: 1.7063
Epoch [158/200], Loss: 1.9796
Epoch [159/200], Loss: 1.9819
Epoch [160/200], Loss: 2.4385
Epoch [161/200], Loss: 1.5695
Epoch [162/200], Loss: 1.6115
Epoch [163/200], Loss: 1.6151
Epoch [164/200], Loss: 2.1459
Epoch [165/200], Loss: 2.3970
Epoch [166/200], Loss: 2.2453
Epoch [167/200], Loss: 1.8896
Epoch [168/200], Loss: 1.8341
Epoch [169/200], Loss: 2.0265
Epoch [170/200], Loss: 2.2931
Epoch [171/200], Loss: 1.7738
Epoch [172/200], Loss: 1.8996
Epoch [173/200], Loss: 2.0559
Epoch [174/200], Loss: 2.3785
Epoch [175/200], Loss: 1.7810
Epoch [176/200], Loss: 1.9395
Epoch [177/200], Loss: 2.1476
Epoch [178/200], Loss: 1.5677
Epoch [179/200], Loss: 1.9600
Epoch [180/200], Loss: 1.8884
Epoch [181/200], Loss: 2.0575
Epoch [182/200], Loss: 1.3027
Epoch [183/200], Loss: 2.1412
Epoch [184/200], Loss: 1.7419
Epoch [185/200], Loss: 1.8414
Epoch [186/200], Loss: 2.4970
Epoch [187/200], Loss: 1.4430

```

```

Epoch [188/200], Loss: 1.9959
Epoch [189/200], Loss: 1.5419
Epoch [190/200], Loss: 2.1794
Epoch [191/200], Loss: 1.6998
Epoch [192/200], Loss: 1.3166
Epoch [193/200], Loss: 1.8938
Epoch [194/200], Loss: 1.6965
Epoch [195/200], Loss: 2.0326
Epoch [196/200], Loss: 1.8422
Epoch [197/200], Loss: 1.9226
Epoch [198/200], Loss: 1.7423
Epoch [199/200], Loss: 1.8963
Epoch [200/200], Loss: 1.5288

```

```
# Prediction 1
```

```
pred_img_path = '/content/drive/My Drive/Co
# image_size = 224 # Resize to this if nec
```

```
# Read and preprocess the image
```

```
pred_img_array = cv2.imread(pred_img_path)
pred_img_array = cv2.resize(pred_img_array,
pred_img_tensor = transforms.ToTensor()(pre
```

```
# Feed the model for prediction
```

```
with torch.no_grad():
    pred_val = model(pred_img_tensor)
    predicted_breed = sorted(new_list)[torc
```

```
from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_p
# Display the predicted breed
print("Predicted Breed for this Dog is:", p
```

```
# Prediction 2
```

```
pred_img_path = '/content/drive/My Drive/Co
# image_size = 224 # Resize to this if nec
```

```
# Read and preprocess the image
```

```
pred_img_array = cv2.imread(pred_img_path)
pred_img_array = cv2.resize(pred_img_array,
pred_img_tensor = transforms.ToTensor()(pre
```

```
# Feed the model for prediction
```

```
with torch.no_grad():
    pred_val = model(pred_img_tensor)
    predicted_breed = sorted(new_list)[torc
```

```

from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_p
# Display the predicted breed
print("Predicted Breed for this Dog is:", p

# Prediction 3
pred_img_path = '/content/drive/My Drive/Co
# image_size = 224 # Resize to this if nec

# Read and preprocess the image
pred_img_array = cv2.imread(pred_img_path)
pred_img_array = cv2.resize(pred_img_array,
pred_img_tensor = transforms.ToTensor()(pre

# Feed the model for prediction
with torch.no_grad():
    pred_val = model(pred_img_tensor)
    predicted_breed = sorted(new_list)[torc

from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_p
# Display the predicted breed
print("Predicted Breed for this Dog is:", p

# Prediction 4
pred_img_path = '/content/drive/My Drive/Co
# image_size = 224 # Resize to this if nec

# Read and preprocess the image
pred_img_array = cv2.imread(pred_img_path)
pred_img_array = cv2.resize(pred_img_array,
pred_img_tensor = transforms.ToTensor()(pre

# Feed the model for prediction
with torch.no_grad():
    pred_val = model(pred_img_tensor)
    predicted_breed = sorted(new_list)[torc

from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_p
# Display the predicted breed
print("Predicted Breed for this Dog is:", p

# Prediction 5

```



```

pred_img_path = '/content/drive/My Drive/Co
# image_size = 224 # Resize to this if nec

# Read and preprocess the image
pred_img_array = cv2.imread(pred_img_path)
pred_img_array = cv2.resize(pred_img_array,
pred_img_tensor = transforms.ToTensor()(pre

# Feed the model for prediction
with torch.no_grad():
    pred_val = model(pred_img_tensor)
    predicted_breed = sorted(new_list)[torc

from google.colab.patches import cv2_imshow
cv2_imshow(cv2.resize(cv2.imread(pred_img_p
# Display the predicted breed
print("Predicted Breed for this Dog is:", p


```



Predicted Breed for this Dog is: whipp



Predicted Breed for this Dog is: whipp





Predicted Breed for this Dog is: whipp



Predicted Breed for this Dog is: dingo



Predicted Breed for this Dog is: ibiza

[Change runtime type](#)